



# **Classification with machine learning: part B**

**Dr Maria Anastasiadi**  
(m.anastasiadi@cranfield.ac.uk)

14<sup>th</sup> January 2025

[www.cranfield.ac.uk](http://www.cranfield.ac.uk)

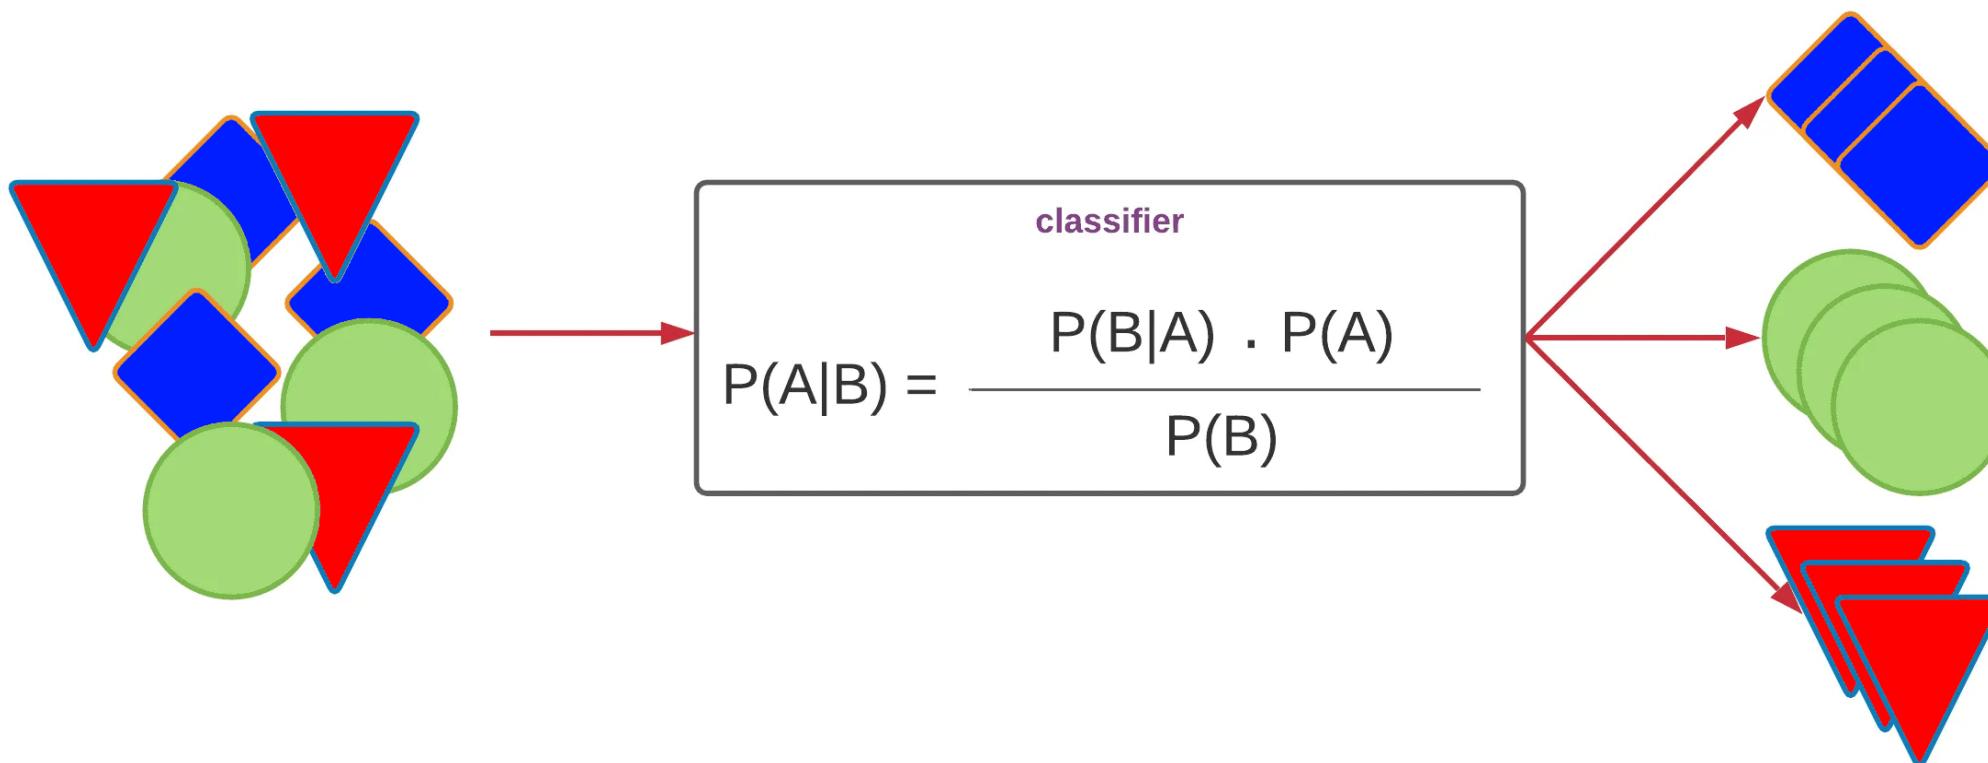


## Intended learning outcomes (ILOs)

- Be able to describe the main principles behind
  - Naïve Bayes
  - SVM
- Provide examples of applications for each algorithm.
- Become familiar with the concept of ML model tuning and optimisation

# Probabilistic learning

## Naive Bayes Classifier



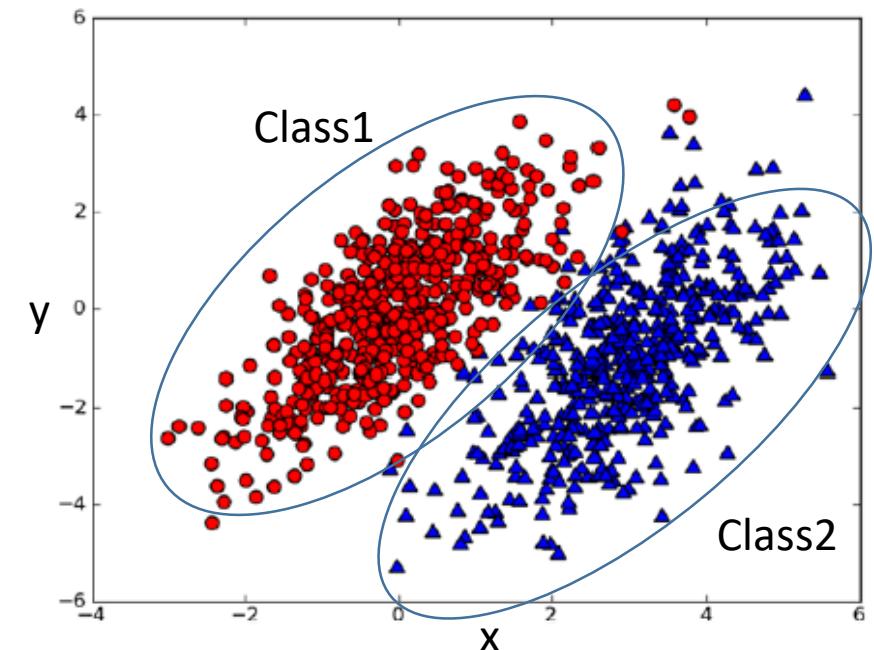
# Classification Using Naïve Bayes

## Naïve Bayes in a nutshell:

1. Assume the data in a dataset belong to two classes Class1 and Class2.
2. Assume we have an equation for the probability of a sample belonging to Class 1  $p_1(x, y)$  and an equation for the probability of a sample belonging to Class 2  $p_2(x, y)$ .
3. To classify a new measurement with features  $(x, y)$  we use the following rules:
  - a) If  $p_1(x, y) > p_2(x, y)$ , then the class is 1.
  - b) If  $p_2(x, y) > p_1(x, y)$ , then the class is 2.

Choose the decision with the highest probability!

**Maximum A Posteriori (MAP)**





# Bayesian probability

## 1. Independent events

The probability of events A and B occurring at the same time:  $P(A \cap B) = P(A) * P(B)$

joint  
probability

## 2. Dependent events

$$P(A \cap B) = P(A) * P(B | A) = P(B) * P(A | B)$$

Conditional probability

## Bayes Theorem:

$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{P(B|A)P(A)}{P(B)}$$

posterior  
probability

likelihood

prior  
probability

marginal probability



# Classifying with conditional probabilities

Using the Bayes theorem, we can calculate the  $p_1$  and  $p_2$  probabilities from the previous example and use it to classify an unseen example.

Given a point identified as  $x,y$ .

What is the probability it belongs to  $c_1$ ?

What is the probability it belongs to  $c_2$ ?

}

$$p(c_1|x \cap y) = \frac{p(x \cap y|c_1)p(c_1)}{p(x \cap y)}$$

$$p(c_2|x \cap y) = \frac{p(x \cap y|c_2)p(c_2)}{p(x \cap y)}$$

As the number of features increases , the formula can get very complicated and computationally demanding.

For example:

$$p(c_1|x \cap y \cap z \cap k) = \frac{p(x \cap y \cap z \cap k|c_1)p(c_1)}{p(x \cap y \cap z \cap k)}$$



# The Naive Bayes algorithm

Naive Bayes assumes **class-conditional independence**, which means that events are independent so long as they are conditioned on the same class level. Also, the denominator can be treated as a constant value. Thus:

$$p(C_1|x \cap y \cap z \cap k) \propto p(x|C_1) p(y|C_1) p(z|C_1) p(k|C_1)$$

The general formula for the Naive Bayes classification algorithm can be summarized by the following equation:

$$P(C_L|F_1, \dots, F_n) = \frac{1}{Z} p(C_L) \prod_{i=1}^n p(F_i|C_L)$$

Where:

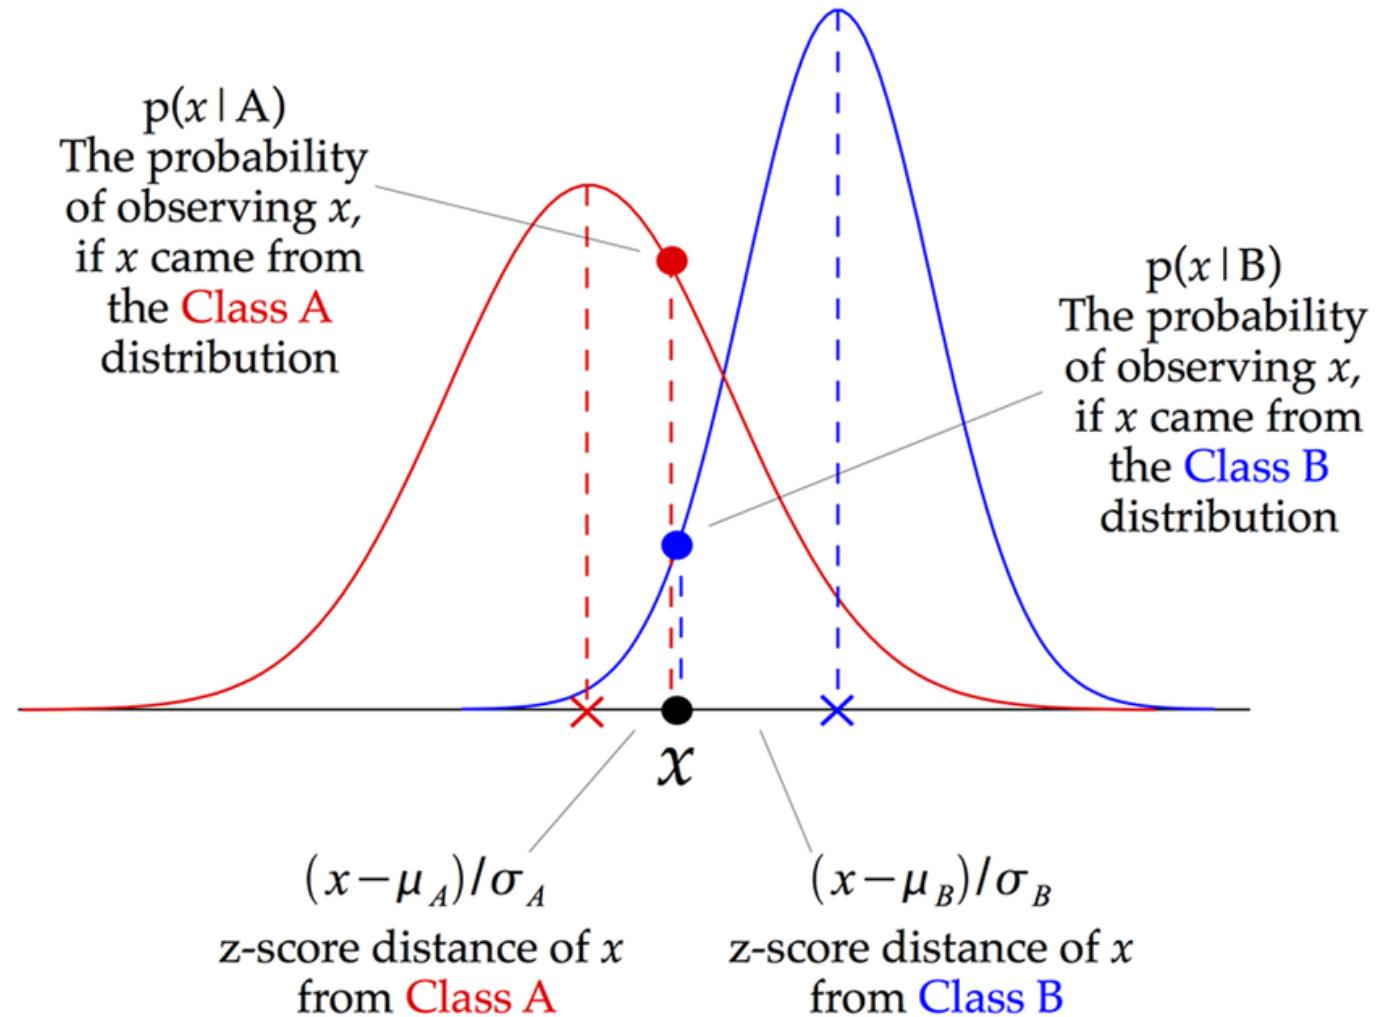
$C_L$  = the level L of class C,

$P(C_L)$  = the prior probability of the class level

$F_i$  = feature i,

$1/Z$  = a scaling factor

# Gaussian Naïve Bayes





# The Naive Bayes algorithm

The Naive Bayes algorithm makes some "naive" assumptions about the data. In particular, Naive Bayes assumes that all of the features in the dataset are equally important and independent, which is rarely true in most real-world applications.

However, in most cases when these assumptions are violated, Naive Bayes still performs fairly well.

There are different types of Naive Bayes Algorithms like GaussianNB, MultinomialNB, BernoulliNB.

Strengths	Weaknesses
Simple, fast, effective	Relies on an often faulty assumption of equally important and independent features
Perform well with noisy and/or missing data	Not ideal for datasets with mainly numeric features
Requires relatively few examples for training	Estimated probabilities are less reliable than the predicted classes
Easy to obtain the estimated probability for a prediction	



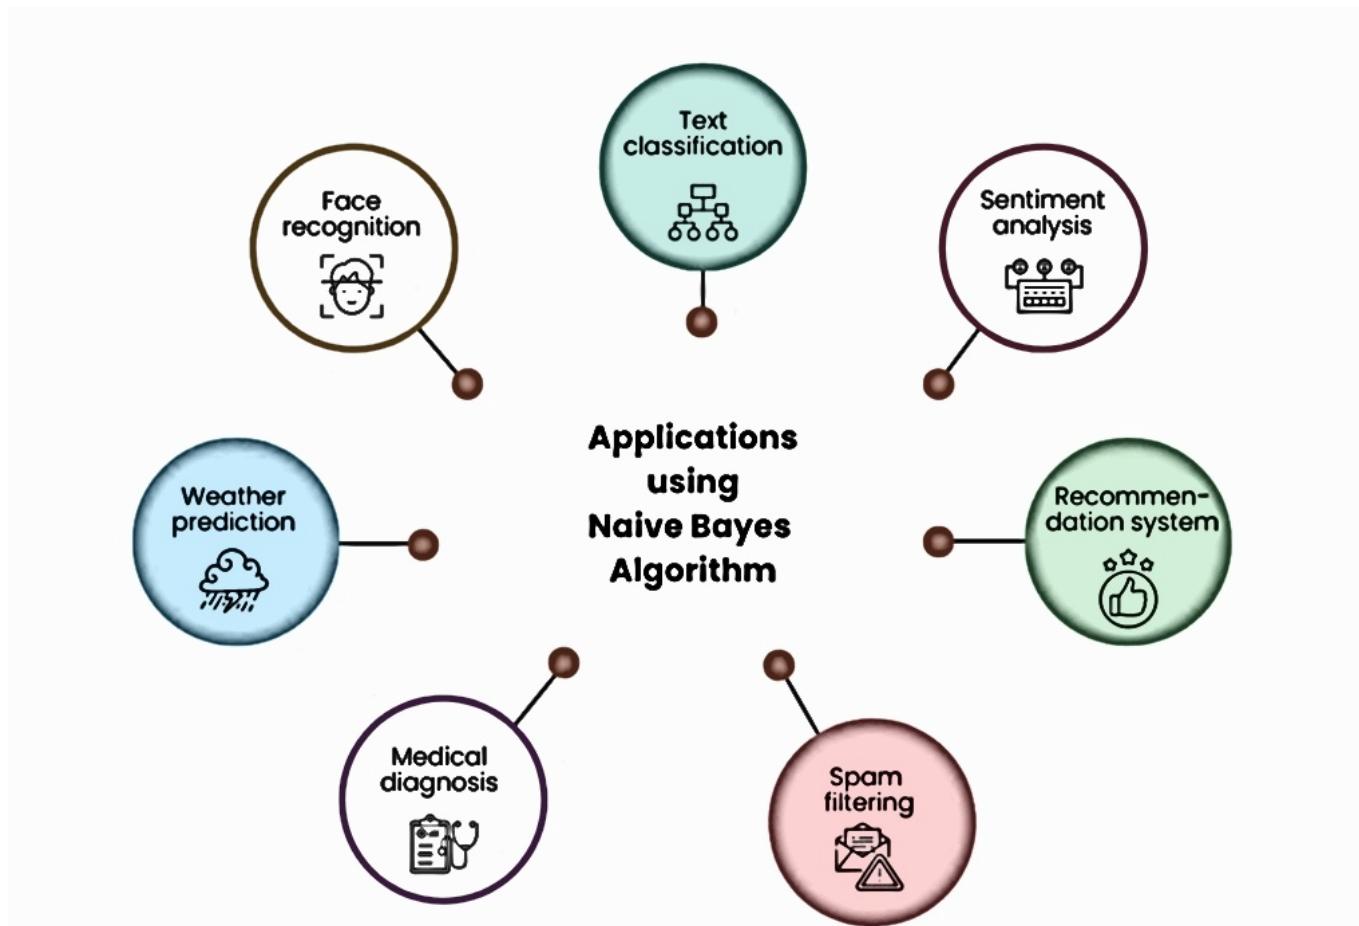
# Naïve Bayes

## How it works

1. Collect data
2. Prepare the data
3. Train a model on the data: Calculate the conditional probabilities of the independent features.
4. Evaluate model performance: Test set. Calculate the error rate.
5. Improve model performance.

**Example:** Classifying text messages (spam / non-spam) – A classical application of NB.

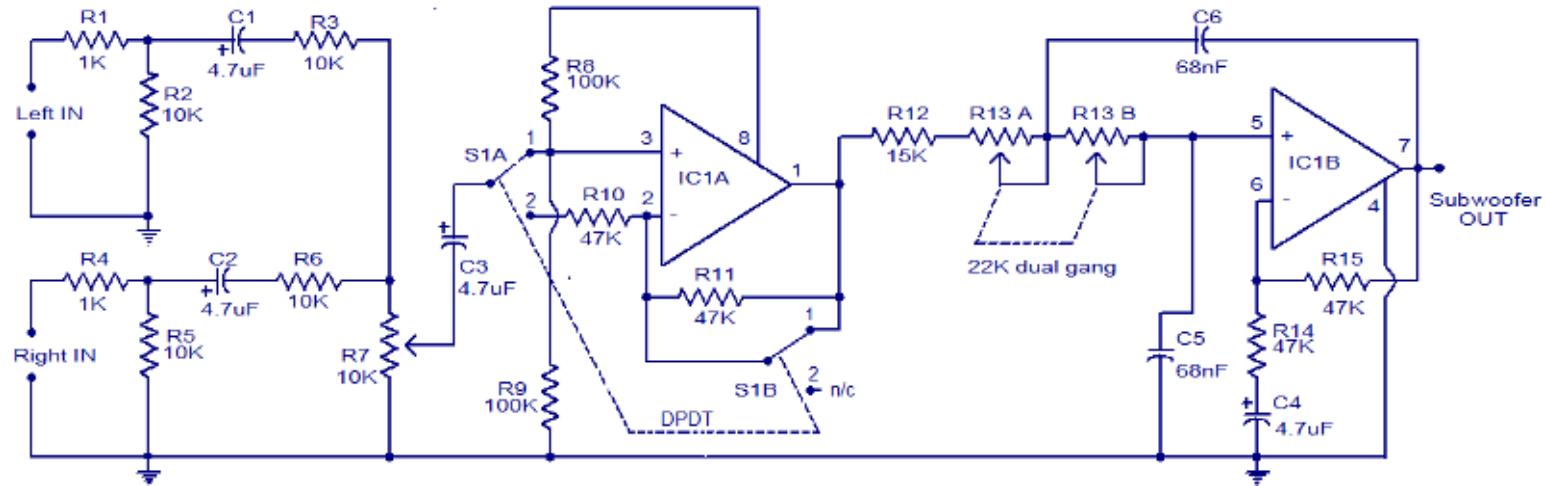
# Naïve Bayes: Applications



# ML: Black Box methods

## The black box metaphor

Imagine a box with a set of inputs (switches/buttons) and a set of outputs (lights)



Can you explain **how** the box works?  
 Can you explain **why** each component  
 has been placed there?

complexity arises from the interaction of  
 many simple components

# ML: Black Box methods

In the case of machine learning, the ‘black box’ metaphor is used for some methods due to the complex mathematics allowing them to function.



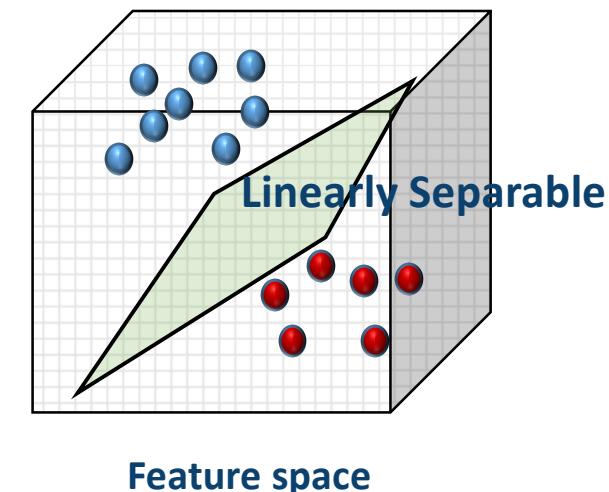
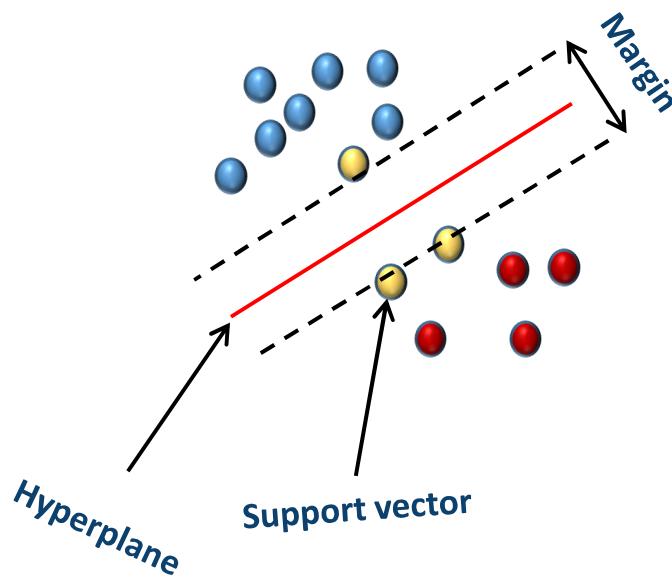
**For example:**

- 1) Neural networks** mimic the structure of animal brains to model arbitrary functions.
- 2) Support vector machines** use multidimensional surfaces to define the relationship between features and outcomes.

# ML: Support Vector Machines (SVM)

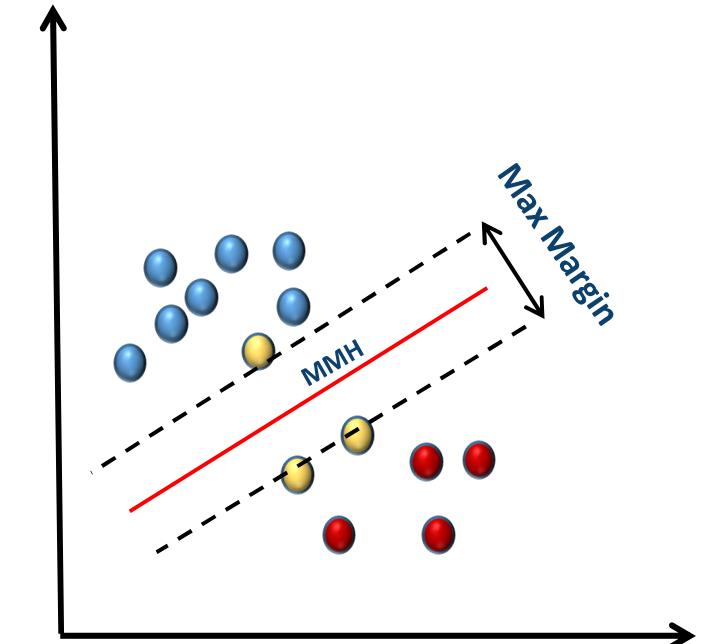
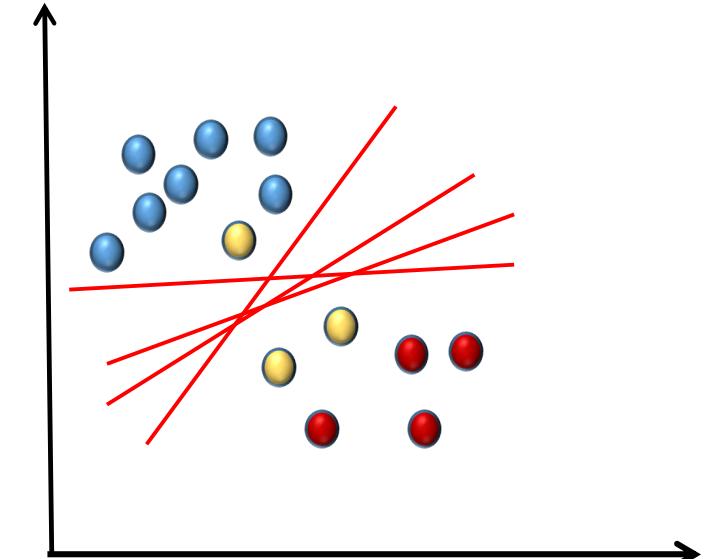
**SVM aims** to create a **linear decision surface (hyperplane)** that can separate classes and has the largest distance (margin) between border-line samples (“**support vectors**”).

SVMs have traditionally been applied for binary classification.



# ML: Support Vector Machines (SVM)

- An infinite number of hyperplanes may exist.
- How does the algorithm choose?
- Search for the **Maximum Margin Hyperplane (MMH)** that creates the greatest separation between the two classes.
- **Important!** Each class must have at least one support vector.
- Using the support vectors alone, it is possible to define the MMH.



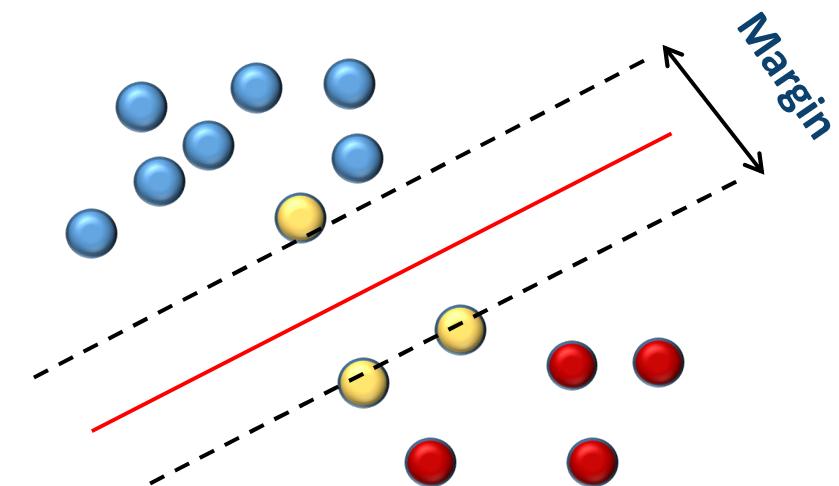
# ML: Support Vector Machines (SVM)

## Methods to find MMH

### Case1: Linearly separable data

Strategy:

- a) Get your dataset
- b) Find a hyperplane which separates the data linearly.
- c) Select two hyperplanes which separate the data with no points between them.
- d) Maximise their distance.





# ML: Support Vector Machines (SVM)

## Step 1: Get your dataset.

Most of the time your data will be composed of  $n$  vectors  $x_i$  containing the variable values for each samples.

Each  $x_i$  will also be associated with a value  $y_i$  indicating if the element belongs to the class (+1) or not (-1).

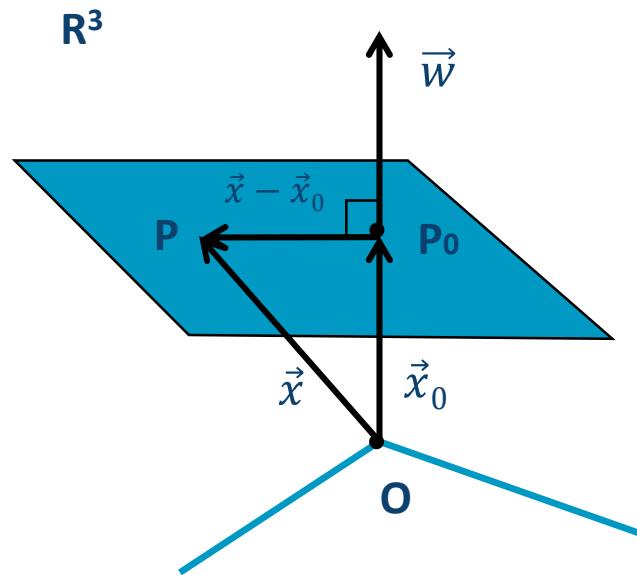
Remember that  $y_i$  can have two possible values -1 or +1.

The initial dataset is defined as :

$$D = \{(x_i, y_i) | x_i \in \mathbb{R}^p, y_i \in \{-1, 1\}\}, \quad i = 1, \dots, n$$

# ML: Support Vector Machines (SVM)

## Step 2. Find the equation for the hyperplane linearly separating the data.



An equation of a hyperplane ( $H_0$ ) is defined by a point ( $P_0$ ) and a vector ( $\vec{w}$ ) perpendicular to the plane.

Let  $P$  be an arbitrary point on a hyperplane.

Then  $\vec{x}_0 = \overrightarrow{OP_0}$  and  $\vec{x} = \overrightarrow{OP}$

A condition for  $P$  to be on the plane is that the vector  $\vec{x} - \vec{x}_0$  is perpendicular to  $\vec{w}$ :

$$\vec{w} \cdot (\vec{x} - \vec{x}_0) = \mathbf{0} \text{ or}$$

$$\vec{w} \cdot \vec{x} - \vec{w} \cdot \vec{x}_0 = \mathbf{0} \quad \text{Let us define } -\vec{w} \cdot \vec{x}_0 = b$$

$$\boxed{\vec{w} \cdot \vec{x} + b = \mathbf{0}}$$

$\vec{w}$  is a vector of  $n$  weights  
And  $b$  is a number known as bias

Holds for  $R^n$  when  $n > 3$



# ML: Support Vector Machines (SVM)

**Step 3. Select two hyperplanes which separate the data with no points between them.**

We have already found a hyperplane  $H_0$  separating the dataset and satisfying  
 $\vec{w} \cdot \vec{x} + b = 0$ .

Now, let's define two other hyperplanes  $H_1$  and  $H_2$  which also separate the data and have the following equations:

$$\vec{w} \cdot \vec{x} + b = +1$$

$$\vec{w} \cdot \vec{x} + b = -1$$

The selected hyperplanes must meet the following constraints:

For each vector  $x_i$  either:

$$\vec{w} \cdot \vec{x} + b \geq +1, \text{ for } x_i \text{ belonging to class } +1$$

or

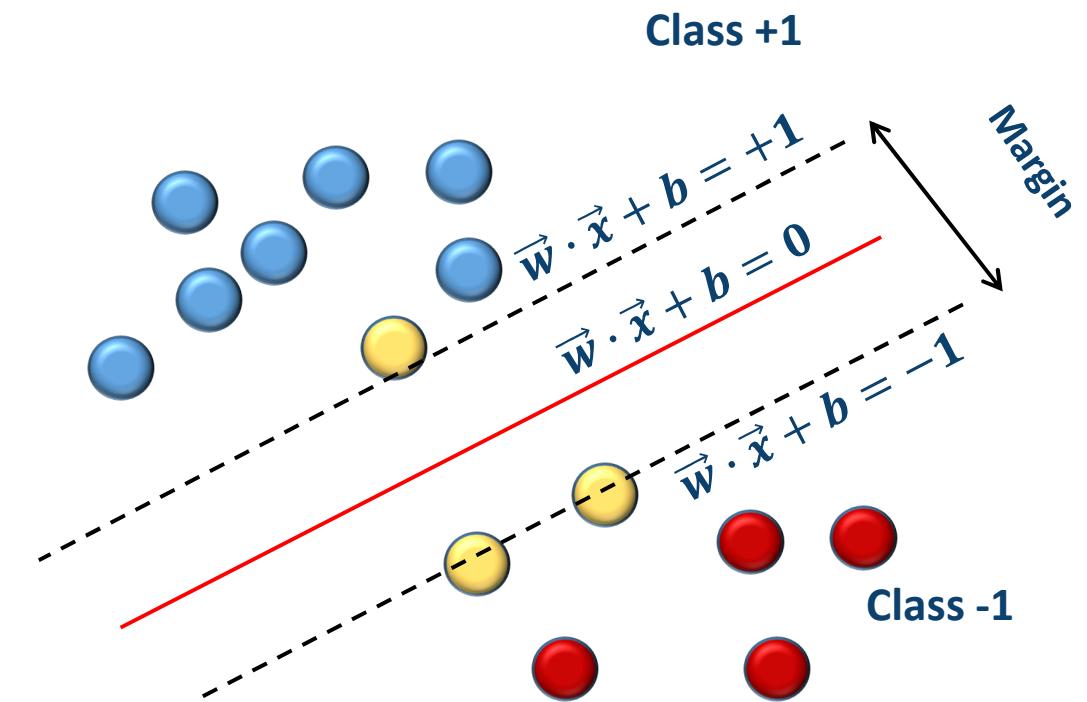
$$\vec{w} \cdot \vec{x} + b \leq -1, \text{ for } x_i \text{ belonging to class } -1$$

# ML: Support Vector Machines (SVM)

If we multiply both sides of the previous constraints with  $y_i$  (+1 for class 1 and -1 for class -1) they can be combined in one equivalent constraint:

$$y_i(\vec{w} \cdot \vec{x} + b) \geq 1$$

This constraint also makes sure there will be no points between the two hyperplanes.





# Support Vector Machines (SVM)

## Step 4. Maximise the distance between H1 and H2.

The distance between these two planes is:  $D = \frac{2}{\|\vec{w}\|}$ ,

where  $\|\vec{w}\|$  is the **Euclidean norm** (the distance from the origin to vector w).

**Therefore, to maximise distance we need to minimise  $\|\vec{w}\|$ .**

The task is typically re-expressed as a set of constraints, as follows:

$$\min \frac{1}{2} \|\vec{w}\|^2$$

Subject to constraint:  $y_i(\vec{w} \cdot \vec{x}_i + b) \geq 1$ , for any  $i = (1, \dots, n)$

# ML: Support Vector Machines (SVM)

## Case 2: Non-linearly separable data – Soft Margins

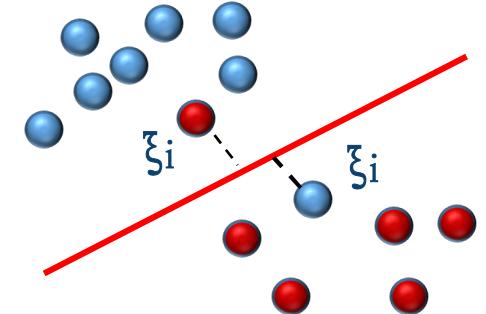
If our dataset isn't linearly separable we can apply **soft margin SVM**: we allow some samples to stray over the line into the margin and be misclassified.

This is done using the **slack variable ( $\xi$ )**. A **cost value (C)** is applied to all points that violate the constraints, and rather than finding the maximum margin, the algorithm attempts to **minimise the total cost**.

$$\min \frac{1}{2} \|\vec{w}\|^2 + C \sum_{i=1}^n \xi_i$$

Subject to constraint:

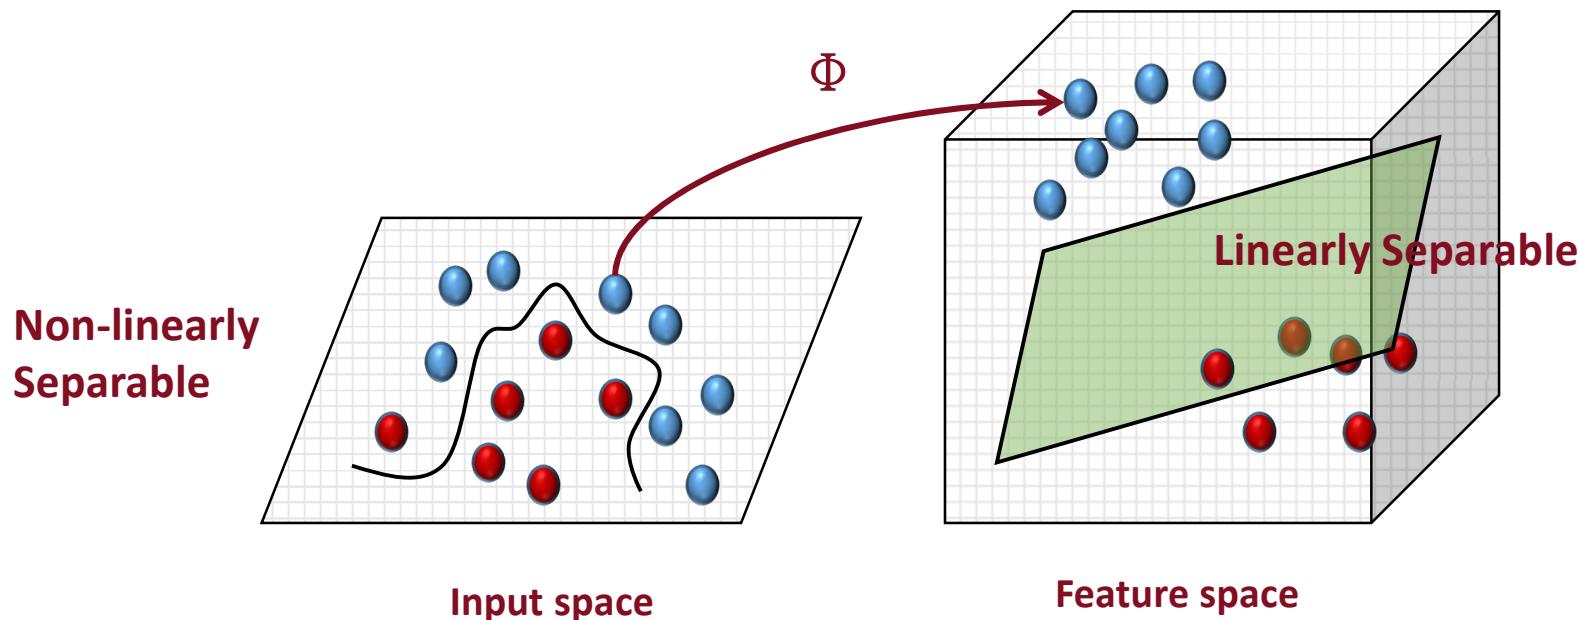
$$y_i(\vec{w} \cdot \vec{x}_i + b) \geq 1 - \xi_i, \text{ for any } i = (1, \dots, n), \xi_i \geq 0$$



# ML: Support Vector Machines (SVM)

## Case 2: Non-linearly separable data - Kernels

When separation is non-linear, the data is plotted into a higher dimensional feature space (“hypercube”) using a kernel function (the kernel trick).





# ML: Support Vector Machines (SVM)

## Non-linearly separable data

### Kernel functions

The kernel “trick” is used to apply some transformation to the feature vectors  $x_i$  and  $x_j$  and combines them using the dot product, which takes two vectors and returns a single number.

### General kernel function

$$K(\vec{x}_i, \vec{x}_j) = \varphi(\vec{x}_i) \cdot \varphi(\vec{x}_j)$$



# ML: Support Vector Machines (SVM)

**Non-linearly separable data**

**Kernel functions**

1. **The linear kernel.** Equivalent to no kernel.

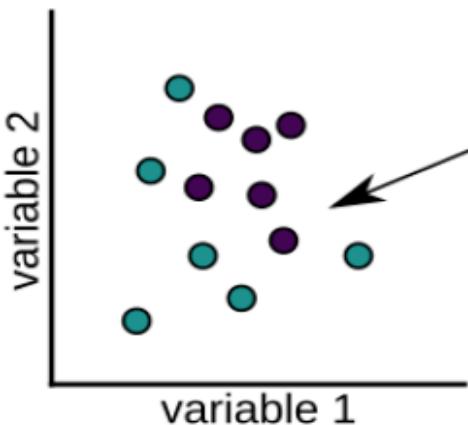
$$K(\vec{x}_i, \vec{x}_j) = \vec{x}_i \cdot \vec{x}_j$$

2. **The polynomial kernel of degree  $d$ .** It adds a simple nonlinear transformation of the data.

$$K(\vec{x}_i, \vec{x}_j) = (\vec{x}_i \cdot \vec{x}_j + 1)^d$$

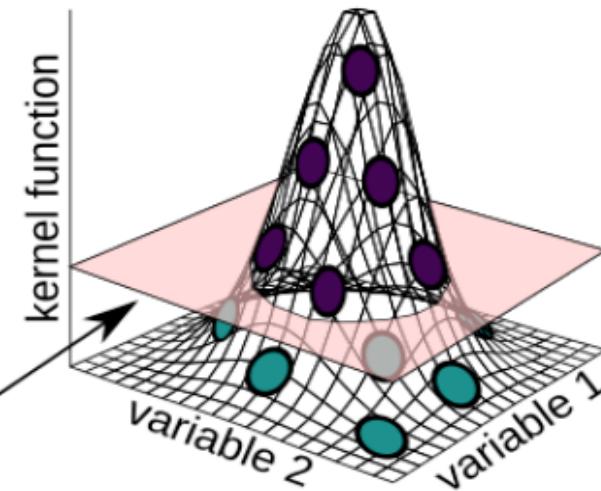
# ML: Support Vector Machines (SVM)

Original data



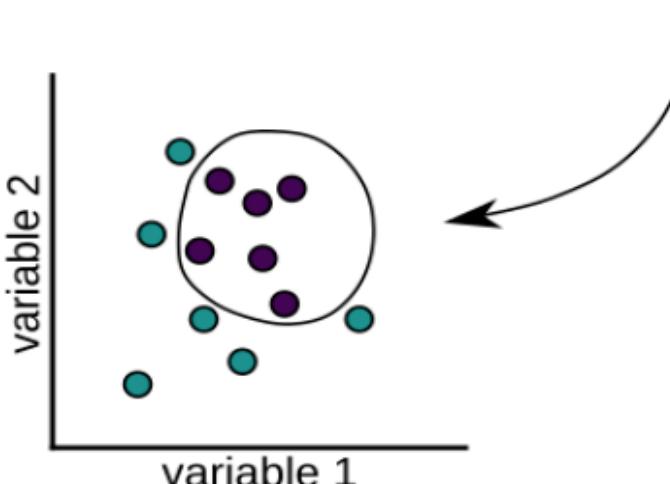
No good linear  
decision boundary

Data with kernel



Linear decision  
hyperplane

polynomial kernel





# ML: Support Vector Machines (SVM)

## Non-linearly separable data

### Kernel functions

**3. The Sigmoid kernel.** Results in a SVM model somewhat analogous to a neural network using a sigmoid activation function.

The Greek letters **kappa** and **delta** are used as kernel parameters:

$$K(\vec{x}_i, \vec{x}_j) = \tanh(k\vec{x}_i \cdot \vec{x}_j - \delta)$$



# ML: Support Vector Machines (SVM)

## Non-linearly separable data

### Kernel functions

#### 4. The Gaussian Radial Basis Function (RBF) kernel.

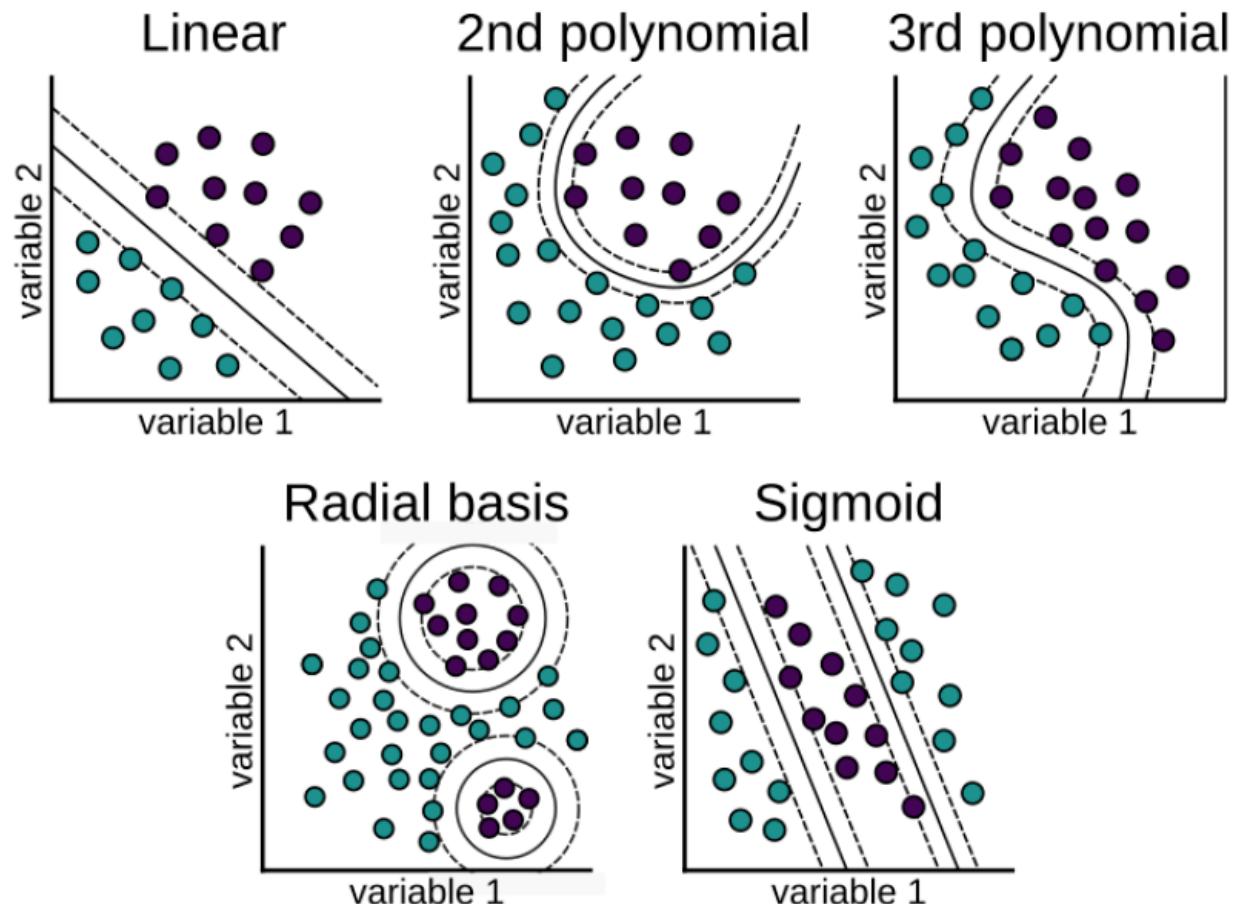
$$K(\vec{x}_i, \vec{x}_j) = e^{\frac{-\|\vec{x}_i - \vec{x}_j\|^2}{2\sigma^2}}$$

- Requires the optimisation of two parameters: **cost** and **gamma**.
- Gamma tends to be the reciprocal of the number of columns in the dataset, or can be varied (e.g. between  $2^{-4}$  and 2).

# ML: Support Vector Machines (SVM)

## Kernel functions

For each Kernel type, the solid line indicates the decision boundary projected back onto the original feature space.



Schematic: Hefin Rhys (2019)



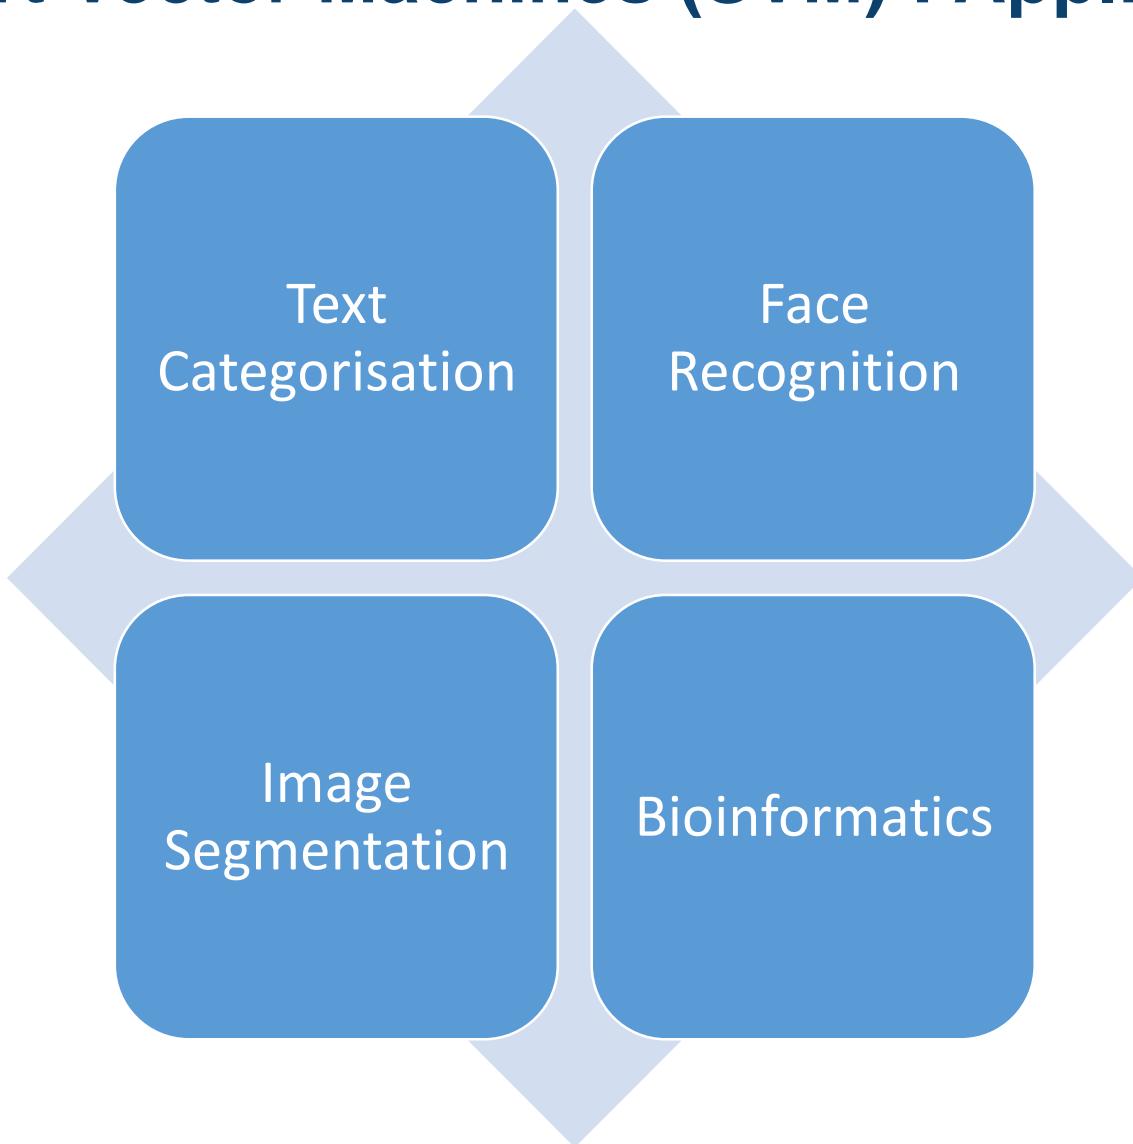
# ML: Support Vector Machines (SVM)

## Kernels pros & cons

Strengths	Weaknesses
Dual use (classification & regression)	Can take time to train for large datasets
Not very prone to overfitting	Various combinations of kernels and model parameters need to be tested
Can cope well with noisy data	Results very difficult to interpret
Widely applied compared to NN	



# ML: Support Vector Machines (SVM) : Applications





# ML: Model Optimisation

## Model Hyperparameters

*Hyperparameters are all the parameters in a ML model that the user defines their values or choose their configuration before training begins.*

E.g. The number of branches in a decision tree, the number of nodes and layers in a NN, the train-test split ratio.

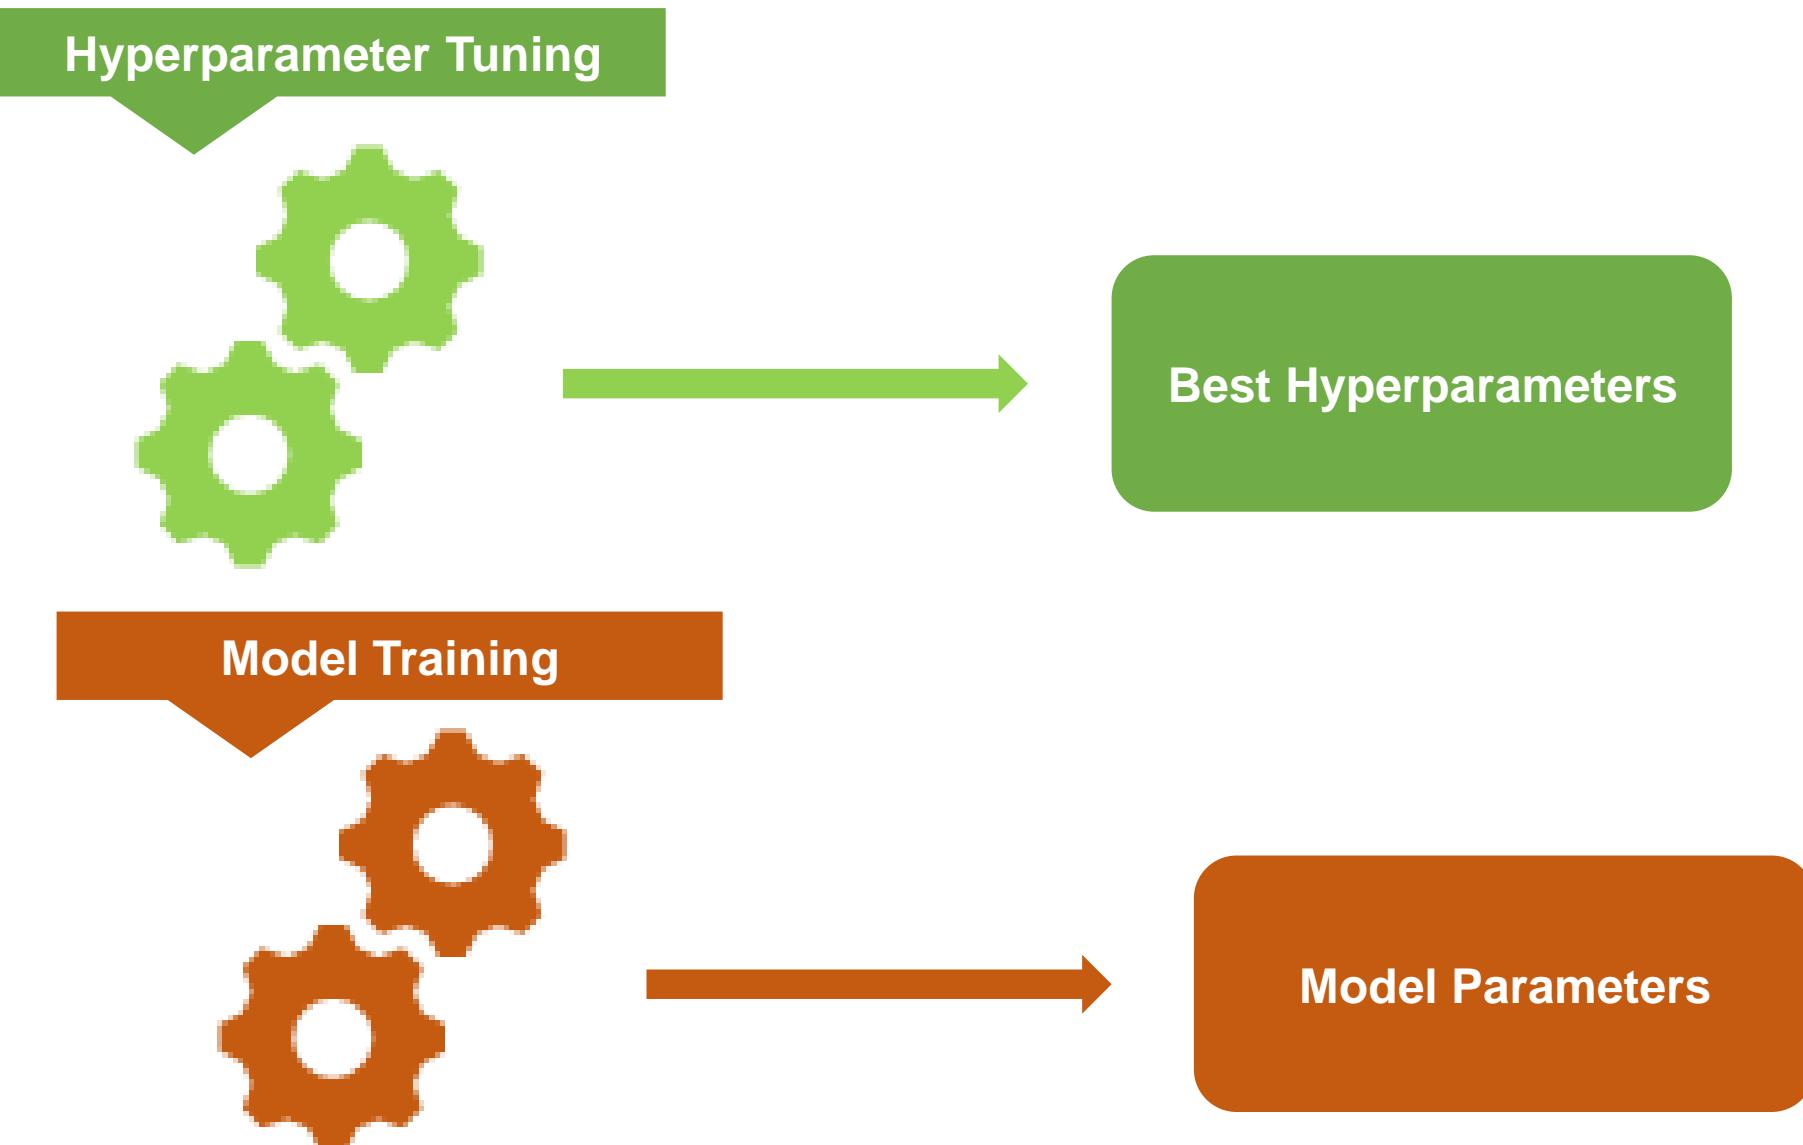
## Model Parameters

*Parameters are internal to the model. They are either learned or estimated purely from the data during training as the algorithm used tries to learn the mapping between the input features and the labels or targets.*

E.g. The coefficients (or weights) of linear and logistic regression models, weights and biases of a NN.



# ML: Model Optimisation





# ML: Hyperparameter Tuning

## Automated hyperparameter tuning

Randomly picking a model's parameter values is a tedious task and not strictly scientific.

- Rather than choosing arbitrary parameter values for each ML model it is better to conduct an automated **grid search** to find the best combination.
- The **caret** package in R provides this option through the **train()** function for over 200 models both for classification and regression tasks.



# ML: Hyperparameter Tuning

**Automated hyperparameter tuning requires the user to consider three questions:**

- 1. What type of ML model should be trained on the data?**
- 2. Which model hyperparameters can be adjusted, and how extensively should they be tuned to find the optimal settings?**
- 3. What criteria should be used to evaluate the models to find the best candidate?**



# ML: Hyperparameter Tuning

Model	Method name	Hyperparameters
k-NN	knn	k
Naïve Bayes	nb	fL, usekernel, adjust
Decision Trees	C5 . 0	model, trials, winnow
OneR Rule Learner	OneR	None
RIPPER Rule Learner	JRip	NumOpt, NumFolds, MinWeights
Linear Regression	lm	None
Regression Trees	rpart	cp
Model Trees	M5	Pruned, smoothed, rules
Neural Networks	nnet	size, decay
Support Vector Machines (Linear Kernel)	svmLinear	C
Support Vector Machines (Radial Basis Kernel)	svmRadial	C, sigma
Random Forests	rf	mtry



# ML: Hyperparameter Tuning

- The goal of **automatic tuning** is to search a set of candidate models comprising a matrix, or **grid**, of parameter combinations.
- By default, **caret** searches at most three values for each of the  $p$  parameters → max of  $3^p$  models.
- For instance, the automatic tuning for **k-NN** will compare 3 models with k=5, k=7 and k=9 (default values).
- If the default grid is not suitable, **caret** allows the user to provide a custom search grid and include more combinations with **expand.grid()**.



# ML: Hyperparameter Tuning

- To see the available parameters for a model we use the function `modelLookup()`.

- For example:

```
> modelLookup("C5.0")
```

	model	parameter	label	forReg	forClass	probModel
1	C5.0	trials # Boosting	Iterations	FALSE	TRUE	TRUE
2	C5.0	model	Model Type	FALSE	TRUE	TRUE
3	C5.0	winnow	Winnow	FALSE	TRUE	TRUE

- The `caret` package also allows the user to calculate a variety of model performance statistics including accuracy and kappa (for classifiers) and R-squared or RMSE (for numeric models).
- Cost-sensitive measures such as sensitivity, specificity, and area under the ROC curve (AUC) can also be extracted.



# ML: Hyperparameter Tuning

**Example of using grid search in a gradient boosting machine (GBM) model.**

There are four main tuning parameters:

- **n.trees** = number of iterations, i.e. trees,
- **interaction.depth** = complexity of the tree (the maximum depth of each tree). Default is 1.
- **shrinkage** = a shrinkage parameter applied to each tree in the expansion (learning rate: how quickly the algorithm adapts)
- **n.minobsinnode** = the minimum number of observations in the terminal nodes of the trees

```
expand.grid(interaction.depth = c(1, 5, 9),  
            n.trees = (1:30)*50,  
            shrinkage = 0.1,  
            n.minobsinnode = 20)
```



# ML: Hyperparameter Tuning

## Example of Random Forests hyperparameter tuning with caret

### 1. Define hyperparameter tuning settings

```
grid <- expand.grid(mtry=2:40)

control <- trainControl(method='repeatedcv',
                        number=10, repeats=5,
                        search='grid')
```

### 2. Train a Random Forests classification model

```
rf.model <- train(Class~, data = TrainSet,
                    method = 'rf',
                    metric = 'Accuracy',
                    ntrees = 800,
                    tuneGrid = grid,
                    trControl = control )
```



# Summary

- Introduction to Probabilistic Methods: Naïve Bayes
- Introduction to Black Box Methods: SVM
- Introduction to hyperparameter tuning concept



**www.cranfield.ac.uk**

**T: +44 (0)1234 750111**

 @cranfielduni

 @cranfielduni

 /cranfielduni